

Grid Features Defense Packet -- XB2IQX

Every trace in this packet is SIMULATED output from a fitted conductance-based single-compartment cell model (src/singlecell_model_v1.py), integrated deterministically at a fixed timestep -- not a raw current-clamp recording. What's being defended here is the feature-extraction methodology applied on top of that model (spike detection, burst/tonic classification, rebound detection, sag depth, adaptation ratio) -- i.e. do these algorithms correctly characterize the dynamics the model actually produces. Every panel below reruns the exact production function on a real simulated trace and marks what it found, so the classification isn't just a text label to take on faith.

cell_floor = -4.50 nA | n grid points = 2500 | burstiness index = 0.042 | F-I slope = 19.91 Hz/nA ($R^2=0.976$)

Contents:

1. Representative traces

One trace per distinct pattern/rebound combination actually observed in this cell's grid.

2. Spike detection

Every detected peak overlaid on real traces.

3. Prominence-threshold sensitivity

How stable spike counts are across a range of thresholds.

4. dV/dt cross-validation

Independent detector cross-check across a random sample.

5. Tonic/bursting/silent classification

ISI sequence + log-ISI KDE evidence.

6. Rebound detection

Recovery-window rebound spikes and the latency cutoff.

7. Sag depth

Baseline, search window, and trough on a real trace.

8. Adaptation ratio

First/last ISI windows on a real, sustained tonic trace.

9. Firing rate / F-I slope

Rate vs. injected current with the fit line.


10. Self-consistency check

Resimulated vs. cached classification match rate.

11. Burst-onset-then-silence detection

Onset-burst detection and the adaptation-ratio ceased-firing gate, contrasted against a burst-onset that keeps firing.


1. Representative traces




One representative trace per distinct (test-window pattern, rebound pattern) combination actually present across this cell's full 2500-point held x injected grid; gray = held baseline, red = test window (injected current), blue = recovery window (released back to held). Chosen by the same selection `plot_example_traces.py` uses for its own per-combination figures (most-prominent point, see that script's `_prominence_score`) -- this is what the cell actually does at each combination, not a hand-picked illustration. Every panel shares the same V (-75 to 18 mV) and time (0-5500 ms) axes, computed from the actual data, so amplitude and duration are directly comparable across panels by eye -- not independently auto-scaled, which would make a subthreshold wobble look as prominent as a real spike train.

`python src/generate_defense_packet.py --cells XB2IQX`
`source: cell_held_injected_grid.pkl -> plot_example_traces.select_exemplar_points() + resimulate_point()`

2. Spike detection




Markers show every peak `scipy.signal.find_peaks` detects at prominence ≥ 15.0 mV (the 15mV absolute floor, since 0.30x this trace's own 47.6mV range would be smaller) -- the exact call `count_spikes_and_rate/compute_isi_ms` use for every spike count, ISI, and burst/tonic classification in the pipeline. 63 spike(s) detected in this window.



3. Prominence threshold sensitivity

XB2IQX -- spike count vs. prominence threshold




PROMINENCE_FRACTION=0.3 is the single most meaningful, least independently-calibrated constant in the pipeline -- every spike count, ISI, and burst/tonic call across the codebase traces back to one `find_peaks(..., prominence=max(range*0.3, 15mV))` call (the absolute floor added 2026-08-13, see section 1's representative traces for why). This sweeps the fraction from 0.10 to 0.55 on real, large-amplitude spiking traces (the floor rarely binds here, since these traces' own range x0.3 already exceeds it) -- a flat plateau around the production value means the spike count this pipeline reports is not sensitive to the exact fraction chosen, not an arbitrary cliff-edge pick.

```
python src/generate_defense_packet.py --cells XB2IQX
source: cell_held_injected_grid.pkl -> resimulate_point() -> scipy.signal.find_peaks swept over PROMINENCE_FRACTION
```


4. dV/dt cross-validation (unused-in-production detector as a sanity check)

example: held=-1.47 inj=-3.70



Cross-check against a second, unused-in-production detector: a candidate peak additionally requires a genuine fast upstroke ($dV/dt \geq 10$ mV/ms within 20 ms before the peak, Bean 2007 convention) before counting. 19 of 19 candidate peak(s) confirmed; the two detectors agree on every spike in this window.



spike count (summed across sample)





Cross-check across a random sample of 20 XB2IQX grid points: every candidate prominence-based peak is additionally required to pass a dV/dt shape-confirmation gate (detect_spikes_dvdt_confirmed -- present in the codebase but not called by the production pipeline). Agreement close to 100% means the simpler production detector isn't missing a shape check that changes the answer.

```
python src/generate_defense_packet.py --cells XB2IQX  
source: cell_held_injected_grid.pkl -> resimulate_point() -> find_silencing_threshold.detect_spikes_dvdt_confirmed()
```


5. Tonic / bursting / silent classification




11 spikes, 10 ISIs. Split at 2 candidate mode(s) -> short
ISI=36.4 ms, long ISI=504.8 ms (ratio 13.88x, needs >= 1.5x). Avg spikes/burst = 2.20 (needs >= 1.5). Final call:
'bursting' (bimodality metric / Ashman's D = 6.33).




6. Rebound (post-inhibitory) detection



1 spike(s) at or after the 5 ms cutoff count as rebound. rebound_occurred=True.




7. Sag depth



Sag depth = baseline - trough = $-43.7 - (-68.1) = 24.4$ mV. Baseline is the held-current settle's final voltage; trough is the minimum voltage over the full 500 ms window (no spike before it).


8. Adaptation ratio



Adaptation ratio = $\text{mean}(\text{last 3 ISIs}) / \text{mean}(\text{first 3 ISIs}) = 9.51$. Firing slows (adapts) over the course of this test window.

9. Firing rate / F-I slope


XB2IQX -- F-I relationship along held=0



39 non-silent, confidently-classified points along the held=0 column (compute_fi_slope's own filter -- flat zero-rate points below rheobase-equivalent excluded so they don't bias the linear fit). This scalar (fi_slope_hz_per_nA) already feeds the cross-cell feature table; this is simply the first time it's been plotted against the data it was fit to.

python src/generate_defense_packet.py --cells XB2IQX
source: cell_held_injected_grid.pkl -> cell_grid_features.pkl (extract_grid_features.compute_fi_slope())

10. Self-consistency check





Re-simulates 40 random cached grid points end-to-end and checks whether the resulting test_pattern/rebound_pattern still matches what's stored in cell_held_injected_grid.pkl. Match rate: 100%. Every held level in the current uniform-grid sweep (except held=0, which reuses the cell's own cached limit-cycle state exactly) is settled via a warm-start from whichever already-settled held level is numerically nearest -- the module docstring in run_held_injected_grid.py already documents that this model shows path-dependent (hysteretic) settling near dynamical transitions, so a resimulation that warm-starts from a different already-settled state than the original sweep used can land on a different classification near a boundary. Any mismatches here are consistent with that

known property, not a new bug.

python src/generate_defense_packet.py --cells XB2IQX
source: cell_held_injected_grid.pkl -> cached test_pattern/rebound_pattern vs. fresh resimulate_point() results

11. Burst-onset-then-silence detection

burst then ceased: held=-2.48 inj=-4.04 -> pattern='tonic', adaptation_ratio=None (suppressed)



Onset burst detected: 4 spikes (mean ISI 14.4 ms) at window start, found by scanning locally for the first ISI jump -- independent of whatever whole-window tonic/bursting/silent label the KDE bimodality test assigns. Trailing silence after the last spike: 2329 ms of a 3000 ms window. likely_ceased_firing=True -- 2329 ms is $\geq 3 \times$ the most recent ISI, so test_adaptation_ratio is suppressed (reported as None): a first-k/last-k ratio isn't a meaningful 'smooth adaptation' number for a train that stopped partway through the window.

burst then sustained: held=-3.77 inj=-3.37 -> pattern='tonic', adaptation_ratio=9.51

